

Interpreting Internal Representations of Transformer Models

Daniel¹² Kiarash¹ Maria³ Vincent¹²

¹Department of Computer Science

²Department of Mathematics

³Department of Mechanical and Industrial Engineering
University of Toronto

May 2025

Table of Contents

- 1 Introduction
- 2 Transformers
- 3 Sparse Autoencoders
- 4 Conclusion

Since 2012, deep learning architectures have achieved notable achievements in the following fields:

- ① Computer vision: object detection, image classification, segmentation
- ② Natural language processing: large language models (ChatGPT, Claude, DeepSeek, etc.), machine translation, coding, math
- ③ Reinforcement learning: robotics, self-driving cars

and many other applied fields, such as health, engineering, etc.

Introduction

- ① Despite these successes, little is known about *how* deep learning models actually process and represent input data
- ② We will examine how input data changes as it flows through a model, focusing on the setting of natural language processing
- ③ Better interpreting how these models work would potentially help us develop safer, more efficient models, and may yield insights into other fields (as many advancements in deep learning already have) such as linguistics, cognitive science, and philosophy.

Prerequisite Vocabulary: Transformers for Text Processing

- **Tokens:** basic unit of text processed by transformers. Can be words, subwords, or characters.
- **Embeddings:** map tokens to dense vectors in a high-dimensional space. These vectors capture semantic relationships: similar tokens (e.g., "king" and "queen") have similar vectors.
 - Embeddings turn discrete tokens into continuous vectors, allowing mathematical operations (e.g., $\text{king} - \text{man} + \text{woman} = \text{queen}$).
- **Positional encoding:** encoding the position of each token in a sentence with a sinusoidal function. Added to the embeddings. Allows the model to learn relative positions of tokens.

Attention is All You Need

- “Attention is All You Need” [Vaswani et al., 2023] introduced the Transformer architecture.
- Attention is a mathematical function defined as:
$$\text{Attention}(Q, K, V) = \text{Softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right) V, \text{ where}$$
 - Q : Queries matrix. Identifies what a token seeks.
 - K : Keys matrix. Identifies what a token offers.
 - V : Values matrix. Carries the actual information to aggregate based on attention weights
 - d_k : Dimension of keys (scaling prevents gradient saturation)
- Q, K, V : learned projections
- Each query vector q_i computes cosine similarity with all keys k_j , creating attention weights that linearly combine value vectors v_j .

Transformer Architecture




Figure: Transformer architecture.

Attention is All You Need

- Instead of one attention function, transformers use $h = 8$ parallel heads:

$$\begin{aligned}\text{MultiHead}(Q, K, V) &= \text{MultiHead}(XW_i^Q, XW_i^K, XW_i^V) \\ &= \text{Concat}(\text{head}_1, \dots, \text{head}_k)W^0\end{aligned}$$

- Each head applies attention to linearly projected versions of Q, K, V : $\text{head}_i = \text{Attention}(XW_i^Q, XW_i^K, XW_i^V)$, where W_i^Q, W_i^K, W_i^V are learnable projection matrices.
- Each head learns a unique 'lens' to focus on different relationships (e.g., syntax vs. semantics).

Our project: Interpretability

Mechanistic interpretability: understanding the internal workings of AI models by identifying the specific mechanisms driving their behaviour.
Why study internal model representations?

- Trust and safety: Transformers are “black boxes”. Learning how knowledge is represented in them helps build trust.
- Model debugging: Understanding attention patterns reveals how models prioritize features.
- Regulation compliance: Laws (like the EU AI Act) often require explanations for automated decisions. Interpretable transformers enable compliance in sectors like healthcare or finance.

Sparse Autoencoders: Intuitive Introduction

- Transformers are powerful but as previously mentioned, they act like “black boxes” so we need to peek inside to understand them.
- Sparse Autoencoders (SAEs) reveal key features the model uses by focusing on sparse and meaningful patterns in its activations.
- Imagine a night before your exam you’re trying to memorize the whole textbook. You might highlight key sentences and ignore the rest. This is similar to how SAEs work (please don’t be like SAEs in real life, the question you think can’t possibly appear on your exam *may* still appear). I tend to be more like SAEs :)

Mathematical Formulation of SAEs

- An SAE has two parts: an encoder and a decoder.
- Encoder: Maps input $\mathbf{x} \in \mathbb{R}^n$ to a latent space $\mathbf{z} \in \mathbb{R}^m$ (where $m > n$).
- Decoder: Reconstructs $\hat{\mathbf{x}} \in \mathbb{R}^n$ from $\mathbf{z}$.
- Loss function:

$$L = \|\mathbf{x} - \hat{\mathbf{x}}\|_2^2 + \lambda \|\mathbf{z}\|_1$$

- First term: Ensures $\hat{\mathbf{x}}$ matches $\mathbf{x}$ (reconstruction).
- Second term: Enforces sparsity in $\mathbf{z}$ using the L1 norm, controlled by λ .
- **Intuition: The SAE learns to represent the input with a few active features (sparse representation) while still being able to reconstruct the input.**

Metrics for Analyzing SAEs

- **Reconstruction Error:** How well does the SAE rebuild the input? Measured with mean squared error (MSE). This is the first term in the loss function.
- **Sparsity:** How few features are active? Often tracked with the L1 norm or count of non-zero elements. This is the second term in the loss function.
- **Monosemanticity:** Do features represent single concepts? Assessed via correlation or human evaluation. This is a new metric introduced in the paper.
- **Interpretability:** How well can we understand the features? This is the main focus of the paper.
- The paper adds five new metrics for auto-interpretation (will be specified with paper details later).

Examples of Auto-Interpreted Latents

- In language models, SAEs might find features like “dog,” “sarcasm,” or “positive sentiment.”
- In vision models, features could be “red circles” or “edges.”
- Paper: “Rethinking Evaluation of Sparse Autoencoders” [\[Source\]](#)




Figure: Example of SAE-identified features in an auto-interpretability pipeline.

Automatically Interpreting Millions of Features

- Paper Title: “Automatically Interpreting Millions of Features in Large Language Models.” [Paulo et al., 2024]
- Tackles interpreting SAEs at scale with:
 - An automated pipeline to label SAE latents.
 - Five new scoring metrics to judge interpretation quality.
 - Open-source code and 1M+ feature interpretations.
- Figure for the pipeline/results:




Figure: Automated interpretation pipeline. [Source]

Conclusion Summary

- 5 New Scoring Metrics:
 - Detection, Fuzzing, Surprisal, Embedding, Intervention
- Important Results:
 - Fuzzing: 0.69 correlation with human eval
 - Intervention: -0.37 correlation with context scores
 - Layer 32: 40% higher scores than Layer 1
- Cost Analysis:
 - Embedding: \$50/100k latents
 - Simulation: 30x more expensive

Conclusion Takeaways

- Automated interpretations help and enable SAE debugging!
- Output features need causal metrics

SAEs can Interpret Randomly Initialized Transformers

From: [Heap et al., 2025]

References



Heap, T., Lawson, T., Farnik, L., and Aitchison, L. (2025).

Sparse autoencoders can interpret randomly initialized transformers.



Paulo, G., Mallen, A., Juang, C., and Belrose, N. (2024).

Automatically interpreting millions of features in large language models.



Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., and Polosukhin, I. (2023).

Attention is all you need.